



Download DZone's 2019 Scaling DevOps Trend Report to learn how to ensure security as you scale DevOps.

[Download Report](#)

SQL With Visual Studio Code

by Simon Foster MVB · Nov. 08, 17 · Database Zone · Tutorial

Writing SQL queries is typically done with SQL Management Studio (SSMS). However, this tool is a bit of a beast, so let's look at how you could use Visual Studio Code instead.

Visual Studio Code is a free text editor — but it is so much more than *just* a text editor. Let's see how. (By the way, VS Code can be downloaded from [here](#).)

To work with SQL Server, download the MS SQL extension. Press **CTRL + SHIFT + P**, select **Install Extension**, and type **mssql**.

Intellisense in Visual Studio Code is brilliant — better than SSMS. Let's look at how to get it all set up.

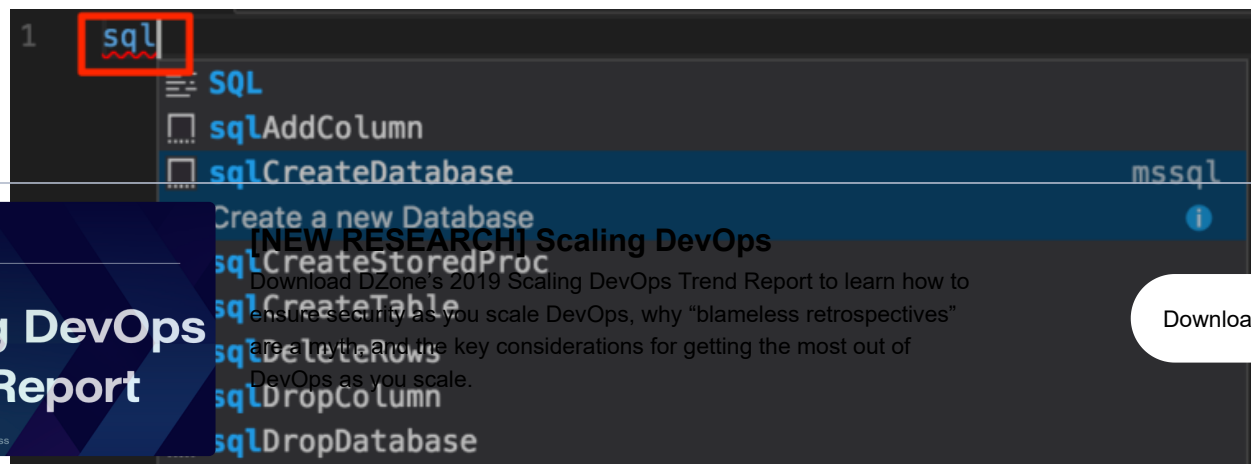
Create a new file and set the language type to SQL (press **CTRL + K + M**).

Open the command palette with **CTRL + SHIFT + P** and type SQL to show the MS SQL commands. Select the **Connect** command.

Then, select **Create Connection Profile**. This creates a profile to connect to your SQL Server. Follow the prompts to get it all set up.

Look in the bottom right corner of the status bar and you should see that you are connected!

Now, if you type **sql**, you will see a long list of SQL code snippets that you can use:



[NEW RESEARCH] Scaling DevOps

Download DZone's 2019 Scaling DevOps Trend Report to learn how to ensure security as you scale DevOps, why "blameless retrospectives" are a myth, and the key considerations for getting the most out of DevOps as you scale.

[Download Now](#)



Choose a snippet to create and edit it as required. When you are happy, press **CTRL + SHIFT + E** to execute.

This is basically all there is to it! However, this is an incredibly powerful way of working; the Intellisense instantly tells you what database objects you can use in your query, and there is a wealth of different snippets you can use.

When returning data, you get a similar view to SSMS, but you can save as Excel, CSV, or JSON.

SSMS is a very graphical way of doing things. You can double-click a table and see its columns or indexes. VS Code relies on T-SQL commands, but you have access to exactly the same information.

For more information about VS Code and the MS SQL extension, check out this documentation.

Like This Article? Read More From DZone



How to Use the SQL Helper Class to Create Web APIs



Database Development in Visual Studio Using SQL Change Automation: Getting Started



How to Hide an Instance of SQL Server



Free DZone Refcard Working With Time Series Data

Topics: DATABASE , SQL , VISUAL STUDIO , MS SQL , SSMS

Published at DZone with permission of Simon Foster, DZone MVB. [See the original article here.](#)

Opinions expressed are solely our own.



On Par With Window Functions
by Keshav M

Scaling DevOps Trend Report
Continued Adoption and Success

[NEW RESEARCH] Scaling DevOps

Download DZone's 2019 Scaling DevOps Trend Report to learn how to manage security in your scaling DevOps, why "blameless retrospectives" are a myth, and the key considerations for getting the most out of DevOps as you scale.

[Download Now](#)

**Use a golf analogy when explaining to executives.
Use a car analogy for all others. — Confucius.**



See window functions more clearly...get it?



[NEW RESEARCH] Scaling DevOps

Download DZone's 2019 Scaling DevOps Trend Report and learn how to ensure security as you scale DevOps, why "blameless retrospectives" are a myth, and the key considerations for getting the most out of DevOps at your scale.

from 40 pages to a few minutes to seconds after using window functions. Query size decreases from 40-pages to a few pages. Back in 2003, every BI layer (like Tableau, and calculating commissions and inventory based on subgroups, positions, etc. These have been in SQL standard in 2003. Every BI layer (like Tableau,

[Download Now](#)

The purpose of analyst efficiency

I've seen real-world running totals,

Looker, Cognos) exploits this functionality.

You may also like: Fun With SQL: Window Functions in Postgres

Introduction to Window Functions

Imagine you have scores of six golfers through two rounds. Now, you need to create the leaderboard and rank them. Rank them using SQL.

Player	Round1	Round2
Marco	75	73
Johan	72	68
Chang	67	76
Isha	74	71
Sitaram	68	72
Bingjie	71	67

Insert the data into Couchbase.

```

1  INSERT INTO golf
2  VALUES("KP1", {"player": "Marco", "round1":75, "round2":73}),
3  VALUES("KP2", {"player": "Johan", "round1":72, "round2":68}),
4  VALUES("KP3", {"player": "Chang", "round1":67, "round2":76}),
5  VALUES("KP4", {"player": "Isha", "round1":74, "round2":71}),
6  VALUES("KP5", {"player": "Sitaram", "round1":68, "round2":72}),
7  VALUES("KP6", {"player": "Bingjie", "round1":71, "round2":67});

```

WITHOUT window functions (current state — Couchbase 6.0)

To write the query, you need a subquery to calculate the rank for each player. This subquery has to scan through all of the data, resulting in the $O(N^2)$, which dramatically increases the execution time and throughput.

```

1  WITH g1 AS (
2  SELECT
3
4
5  (select row 1+COUNT(*)

```



NEW RESEARCH | Scaling DevOps
 Download DZone's 2019 Scaling DevOps Trend Report to learn how to
 ensure security as you scale DevOps, why "blameless retrospectives"
 are a myth, and the key considerations for getting the most out of
 DevOps as you scale.

AS player,
 AS T,
 AS ToPar,

[Download Now](#)

```

5      from g1 as g2
6      where (g2.round1 + g2.round2) <
7             (g3.round1 + g3.round2))[0] AS sqlrankR2
9 FROM g1 as g3
10 ORDER BY sqlrankR2
11
12 Results:
13 T ToPar player sqlrankR2
14 138 -6 "Bingjie"      1
15 140 -4 "Johan"        2
16 140 -4 "Sitaram"      2
17 143 -1 "Chang"        4
18 145 1 "Isha"          5
19 148 4 "Marco"         6

```

With window functions in Mad-Hatter (upcoming release)

This query returns the player, the total after two rounds (T), how much of the score is over/under par (ToPar), and then **rank**s them based on the scores of the first two rounds. This is the NEW functionality in Mad-Hatter. The time complexity of this is $O(N)$, meaning execution time will only increase linearly.

```

1 SELECT  player                AS player,
2          (round1+round2)       AS T,
3          ((round1+round2) - 144) AS ToPar,
4          RANK() OVER(ORDER BY (round1+round2)) AS rankR2
5 FROM golf;
6
7
8 T ToPar player  rankR2
9 138 -6 "Bingjie"  1
10 140 -4 "Johan"    2
11 140 -4 "Sitaram"  2
12 143 -1 "Chang"    4
13 145 1 "Isha"      5
14 148 4 "Marco"     6

```

Observ

The query exp

1. Perform



[NEW RESEARCH] Scaling DevOps

Download DZone's 2019 Scaling DevOps Trend Report to learn how to ensure security as you scale DevOps, why "blameless retrospectives" are a myth, and the key considerations for getting the most out of DevOps as you scale.

[Download Now](#)

scenario is much better. We plan to measure.

2. When the ranking requirements depend on multiple documents, the query becomes quite complex to write, optimize, and run.

3. All of this affects the TCO overall.

Now, let's create an expanded dashboard.

Show add dense rank, row number, who's ahead, and the number of strokes behind the leader. All very common things in a reporting resituation. You're seeing the new window function whenever you see the OVER() clause. The query below has six window functions.

```

1  SELECT      player                      AS player,
2              (round1+round2)            AS T,
3              ((round1+round2) - 144)     AS ToPar,
4              RANK() OVER(ORDER BY (round1+round2)) AS rankR2,
5              DENSE_RANK() OVER (ORDER BY (round1+round2)) AS rankR2Dense,
6              ROW_NUMBER() OVER() rownum,
7              ((round1+round2) -
8                  FIRST_VALUE(round1+round2)
9                      OVER(ORDER BY (round1+round2))) AS strokesbehind,
10             RANK() OVER(ORDER BY (round1))      AS rankR1,
11             LAG(player, 1, "None") OVER(ORDER BY round1+round2)
12                                     AS inFront
13 FROM golf
14 ORDER BY rankR2
15
16
17 T ToPar inFront  player  rankR1  rankR2  rankR2Dense  rownum  strokesbehind
18 138 -6 "None"    "Bingjie" 3 1 1 3 0
19 140 -4 "Johan"   "Sitaram" 2 2 2 2 2
20 140 -4 "Bingjie" "Johan" 4 2 2 4 2
21 143 -1 "Sitaram" "Chang" 1 4 3 1 5
22 145 1 "Chang"    "Isha"   5 5 4 5 7
23 148 4 "Isha"     "Marco" 6 6 5 6 10

```

As you saw earlier, doing this query with **six window functions using** the subquery method will be a larger effort, expensive, error-prone query.

In addition to
each of them a

- RANK()
- DENSE_RANK()
- PERCENT_RANK()
- CUME_DIST()



UNT, SUM, AVG, etc) as window functions. Sitaram has added the following window functions. The syntax and semantic

NEW RESEARCH Scaling DevOps

and well-described in the articles of the reference section below.
Download DZone's 2019 Scaling DevOps Trend Report to learn how to
ensure security as you scale DevOps, why "blameless retrospectives"
are a myth, and the key considerations for getting the most out of
DevOps as you scale.

[Download Now](#)

- [SUM\(\)](#)
- [NTILE\(\)](#)
- [RATIO_TO_REPORT\(\)](#)
- [ROW_NUMBER\(\)](#)
- [LAG\(\)](#)
- [FIRST_VALUE\(\)](#)
- [LAST_VALUE\(\)](#)
- [NTH_VALUE\(\)](#)
- [LEAD\(\)](#)

References

1. Probably the Coolest SQL Feature: Window Functions
2. A Window into the World of Analytic Functions
3. Oracle Reference

Further Reading

- [WINDOW Function](#)
- [How to Avoid Excessive Sorts in Window Functions](#)

Like This Article? Read More From DZone

[Parallelism: Estimated Costs of All Queries](#)

[Keeping MySQL in Check: Real-Life Advice to Tame Daily Performance Degradations](#)

[Verifying Query Performance Using ProxySQL](#)


[Free DZone Refcard
Working With Time Series Data](#)

Topics: DATABASES



Published at DZone
Opinions expressed are

[NEW RESEARCH] Scaling DevOps

WINDOW FUNCTIONS TUTORIAL, CODE SNIPPET, QUERY PERFORMANCE
Download DZone's 2019 Scaling DevOps Trend Report to learn how to ensure security as you scale DevOps, why “blameless retrospectives” are a myth, and the key considerations for getting the most out of DevOps as you scale. [See the original article here.](#) 

[Download Now](#)



[NEW RESEARCH] Scaling DevOps

Download DZone's 2019 Scaling DevOps Trend Report to learn how to ensure security as you scale DevOps, why "blameless retrospectives" are a myth, and the key considerations for getting the most out of DevOps as you scale.

[Download Now](#)